NF-generated samples versus the exact declared tilt target, per trained variant (synthetic B_toy; NLL annotations are a separate, non-bar diagnostic -- not evidence of sampling agreement)

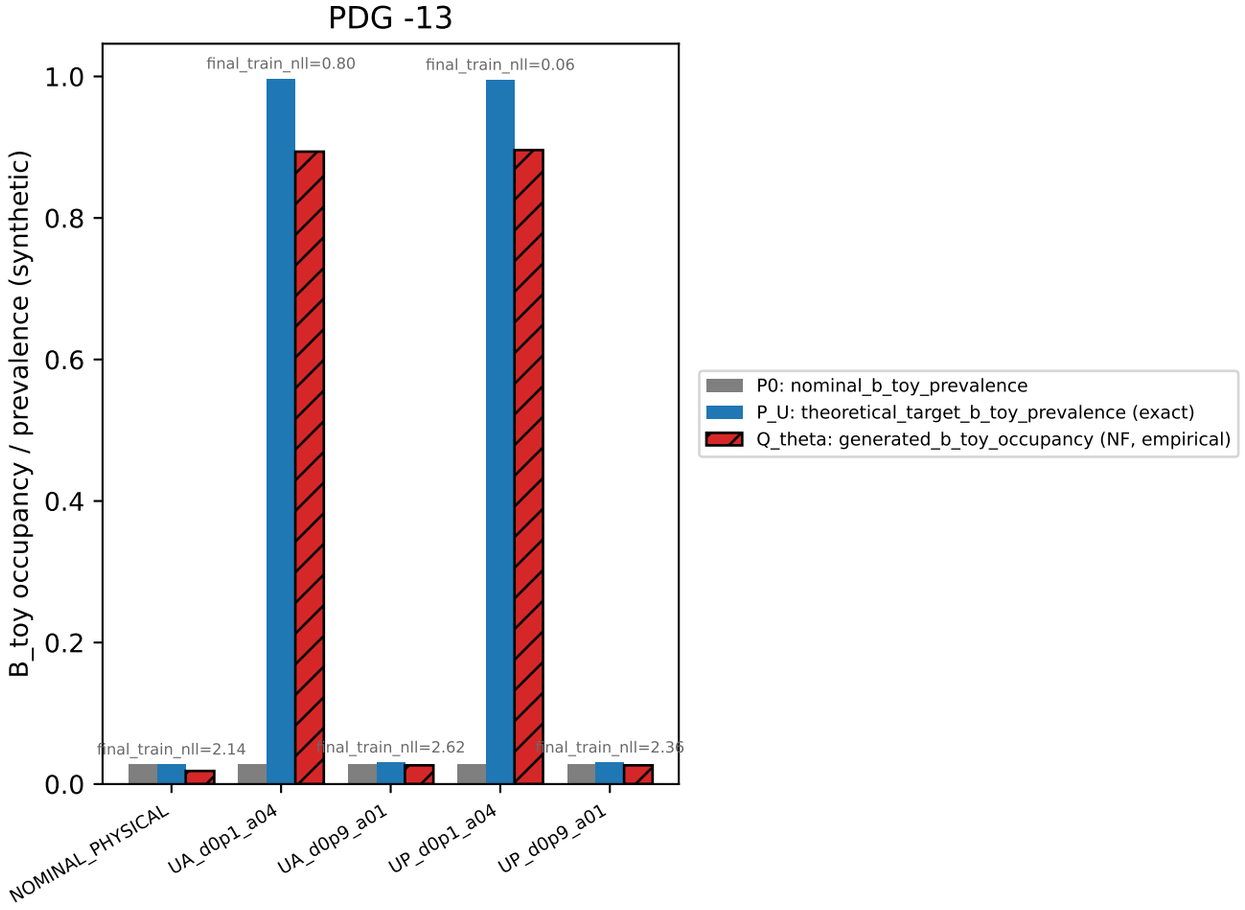

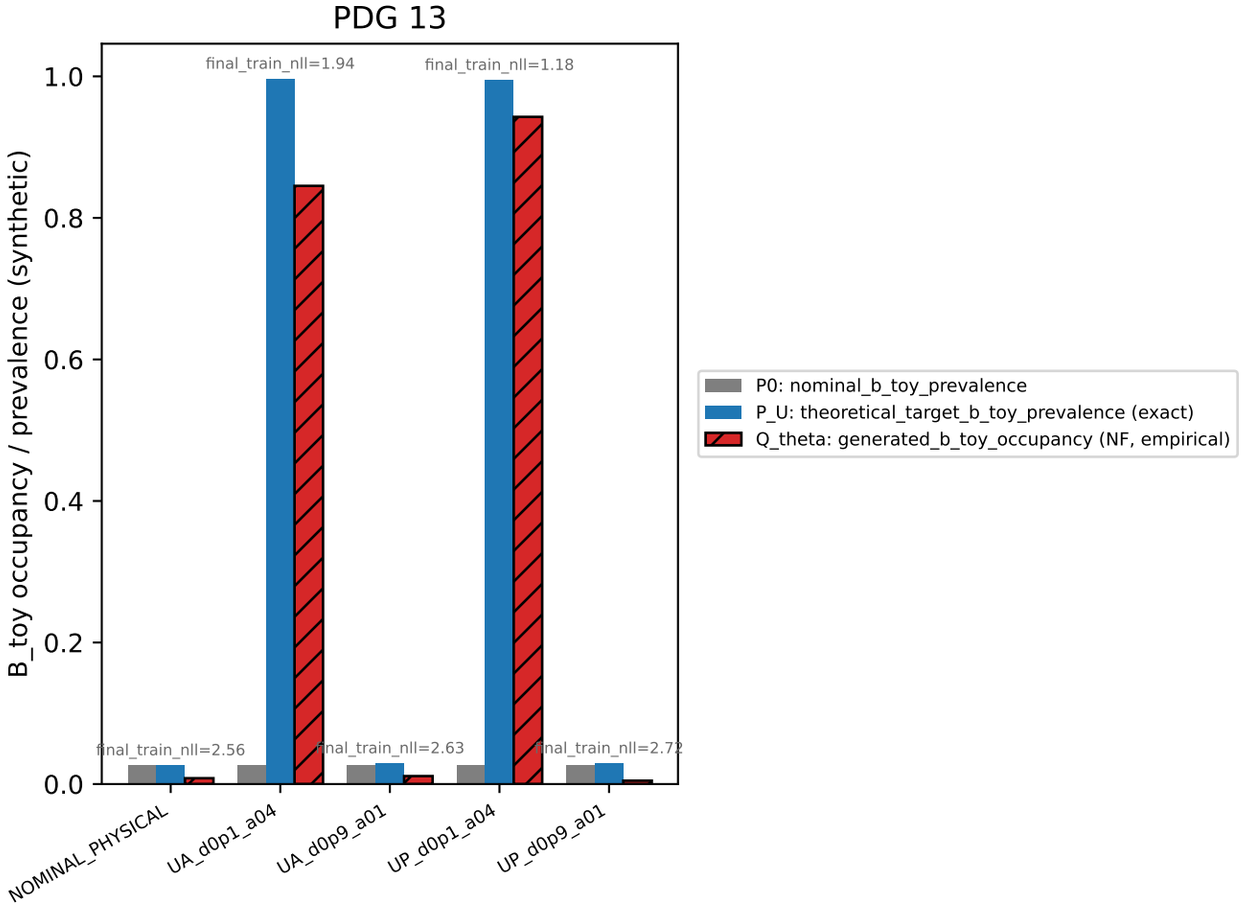

Bounded CPU-scale affine-coupling flow (8 blocks, width 96, <=40 epochs, early stopping), not a converged production model. generated_b_toy_occupancy is measured from samples drawn from the trained flow and mapped back to physical space; it approaches but does not exactly reach the declared theoretical_target_b_toy_prevalence, which is expected at this training budget and is not evidence of a sampler or mathematics bug (those are validated separately and exactly via the closed-form/direct-summation/alias-sampler path).